

Guía Introductoria: Primeros Pasos en R

José R. Caro Barrera (a partir de Saghir Bashir y Dirk Eddebuettel)

This version was compiled on 5 de abril de 2021

Esta guía es un intento de facilitar la puesta en marcha en R y Rstudio. El contenido es muy práctico, pues se usan ejemplos y conjuntos de datos reales, la idea es que el alumno se familiarice con los conceptos clave de R, a la hora de manipular, visualizar y resumir un conjunto de datos.

R | Estadística | Química | RStudio

1. Prólogo

Esta guía de “Introducción” es una breve aproximación de lo que se puede hacer con R¹. Para aprovecharla al máximo es conveniente su lectura y a la vez, ir haciendo los ejemplos y los ejercicios que aquí se presentan usando RStudio².

Esta guía es una variante original del documento³ pero hace hincapié en el uso de Base R junto con una gestión cuidadosa de las dependencias como se explica a continuación.

Experimentar en Modo Seguro. Lo ideal es experimentar con comandos y opciones, ya que es una parte esencial del proceso de aprendizaje. Las cosas pueden salir “mal” (lo harán), como recibir mensajes de error o eliminar cosas que se hayan creado, pero la mayoría de las situaciones pueden recuperarse (por ejemplo, reiniciando R). Para hacer esto “de forma segura”, comienza con una sesión en R nueva sin ningún otro dato cargado.

2. Introducción

Antes de Comenzar. Es conveniente asegurarse de que:

1. R y RStudio (por ese orden) están instalados.
2. Comprobar que estamos en el directorio de trabajo deseado con `getwd()`.

Comenzando R y RStudio. R arranca automáticamente cuando se abre RStudio (ver Fig. 1). La consola comienza con información sobre el número de versión, la licencia y los colaboradores. La última línea es un mensaje estándar “>” que indica que R está listo y esperando instrucciones para hacer algo.

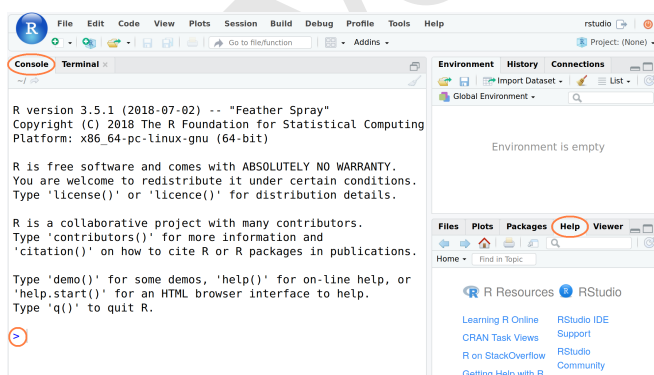


Fig. 1. Captura de pantalla de RStudio con la consola a la izquierda y la pestaña Ayuda en la parte inferior derecha

Saliendo de R y RStudio. Cuando salgamos de RStudio, nos preguntará si deseamos Guardar espacio de trabajo (Save workspace) con dos opciones:

- “Yes” – Nuestro espacio de trabajo actual de R (que contiene el trabajo que hemos realizado) será restaurado la próxima vez que se abra RStudio.
- “No” – Comenzaremos con una nueva sesión de R la próxima vez que abramos RStudio. Por ahora, seleccionaremos “No” para evitar errores de sesiones anteriores.

3. La Ayuda (Help) en R

RStudio tiene un extenso sistema de ayuda incorporado que suele ser muy útil a la hora de encontrar soluciones a cualquier problema ya sea con comandos, paquetes y/o funciones.

Función `help()`. Desde la “consola” de R se puede usar la función `help()` o `?`. Por ejemplo, los dos comandos siguientes proporcionan el mismo resultado:

```
help(mean)
?mean
```

Búsqueda de palabra clave.. Para buscar una palabra clave se usa la función `apropos()` con la palabra en cuestión entre comillas doble (“keyword”) o comillas simple ('keyword'). Por ejemplo:

```
apropos("mean")
# [1] ".colMeans"      ".rowMeans"
# [3] "colMeans"       "frollmean"
# [5] "kmeans"         "mean"
# [7] "mean.cl_boot"   "mean.cl_normal"
# [9] "mean.sdl"       "mean.se"
# [11] "mean.Date"      "mean.default"
# [13] "mean.difftime"  "mean.POSIXct"
# [15] "mean.POSIXlt"   "rowMeans"
# [17] "weighted.mean"
```

Ejemplos de Ayuda. Use la función `example()` para ejecutar los ejemplos al final de la ayuda para una función:

```
example(mean)
#
# mean> x <- c(0:10, 50)
#
# mean> xm <- mean(x)
#
# mean> c(xm, mean(x, trim = 0.10))
# [1] 8.75 5.50
```

La Ayuda de RStudio. Rstudio proporciona un cuadro de búsqueda en la pestaña “Ayuda” para hacer el trabajo más fácil (ver Fig. 1).

¹Proyecto R: <https://www.r-project.org/>

²IDE RStudio: <https://www.rstudio.com/products/RStudio/>

³Getting Started with R: <https://github.com/saghirb/Getting-Started-in-R>

Búsqueda de ayuda para R on-line. Hay muchos recursos on-line que pueden servir de ayuda, sin embargo, copiar y pegar sin más, a veces, no será de mucha ayuda y ralentizará el proceso de aprendizaje y desarrollo. Cuando se haga una búsqueda on-line use [R] en el término de búsqueda, (p. ej. "[R] resumen estadísticos por grupo"). A menudo hay más de una solución para un problema en R, por lo que es bueno investigar las diferentes opciones.

Ejercicio. Intente lo siguiente:

1. `help(median)`
2. `?sd`
3. `?max`

Advertencia. Si un comando R no está completo, R mostrará un signo más (+) en la segunda línea y en las siguientes hasta que la sintaxis del comando sea correcta.

```
+
```

Para cortar esto, presione la tecla de escape (ESC).

Truco. Para recuperar un comando escrito anteriormente, use la tecla de flecha hacia arriba (↑). Para ir entre los comandos escritos anteriormente, use las flechas hacia arriba y hacia abajo (↓). Para modificar o corregir un comando use la flecha izquierda (←) y la derecha (→).

4. Algunos conceptos en R

En lenguaje R, escalares, vectores/variables y datasets se denominan **objetos**. Para crear objetos (cosas) hay que usar el operador asignación `<-`. Por ejemplo, a continuación, al objeto `altura` se le ha asignado el valor 173 (tecleando `altura` se muestra su valor):

```
altura <- 173
altura
# [1] 173
```

Advertencia: sobre las mayúsculas. `edad` y `EdaD` son diferentes:

```
edad <- 10
EdaD <- 50
```

```
edad
# [1] 10
EdaD
# [1] 50
```

Nuevas líneas. Los comandos en R suelen estar separados por una nueva línea, pero también pueden estarlo por un punto y coma ;.

```
Nombre <- "Rafael"; Edad <- 20; Ciudad <- "Madrid"
Nombre; Edad; Ciudad
# [1] "Rafael"
# [1] 20
# [1] "Madrid"
```

Comentarios. Es útil incluir comentarios explicativos en el código. Estos comentarios suelen ser de ayuda en el futuro cuando se retome el script. Los comentarios en R comienzan por el símbolo almohadilla (#). Todo lo que esté después de este símbolo hasta el final de la línea será ignorado por R.

```
# Este comentario será ignorado cuando se ejecute.
Ciudad # El texto despues de "#" es ignorado.
# [1] "Madrid"
```

5. R como calculadora

Es posible usar R como calculadora. Pruebe lo siguiente:

```
2 + 3
# [1] 5
(5*11)/4 - 7
# [1] 6.75
# ^ = "potencia de un número"
7^3
# [1] 343
```

Otras funciones matemáticas. También se pueden utilizar funciones matemáticas estándar que normalmente se encuentran en una calculadora científica.

- Trigonómicas: `sin()`, `cos()`, `tan()`, `acos()`, `asin()`, `atan()`
- Redondeo: `abs()`, `ceiling()`, `floor()`, `round()`, `sign()`, `signif()`, `sqrt()`, `trunc()`
- Logarítmicas y Exponenciales: `exp()`, `log()`, `log10()`, `log2()`

```
# Raíz cuadrada
sqrt(2)
# [1] 1.414214
# Redondeo a la baja al entero más cercano
floor(8.6178)
# [1] 8
# Redondeo a dos decimales
round(8.6178, 2)
# [1] 8.62
```

Ejercicio. Compruebe qué hacen los siguientes pares de comandos:

1. `ceiling(18.33)` y `signif(9488, 2)`
2. `exp(1)` y `log10(1000)`
3. `sign(-2.9)` y `sign(32)`
4. `abs(-27.9)` y `abs(11.9)`

6. Algunos conceptos adicionales en R

Se pueden hacer algunas cosas útiles e interesantes con el operador de asignación `<-`:

```
longitud <- 7.8
anchura <- 6.4
area <- longitud * anchura
area
# [1] 49.92
```

Objetos de texto. También podemos asignar un texto a un objeto.

```
saludo <- "Hola Mundo!"
saludo
# [1] "Hola Mundo!"
```

Vectores. Los objetos presentados hasta ahora han sido todos escalares (valores únicos). Trabajar con vectores es donde R brilla mejor, ya que son los componentes básicos de los conjuntos de datos. Para crear un vector podemos usar la función `c()` (combinar valores en un vector).

```
# Un vector "numérico"
x1 <- c(26, 10, 4, 7, 41, 19)
x1
# [1] 26 10 4 7 41 19
# Un vector caracter de nombres de países
x2 <- c("Brasil", "Italia", "Cuba", "Ghana")
x2
# [1] "Brasil" "Italia" "Cuba" "Ghana"
```

There are many other ways to create vectors, for example, `rep()` (replicate elements) and `seq()` (create sequences):

```
# Repeat vector (2, 6, 7, 4) three times
r1 <- rep(c(2, 6, 7, 4), times=3)
r1
# [1] 2 6 7 4 2 6 7 4 2 6 7 4
# Vector from -2 to 3 incremented by half
s1 <- seq(from=-2, to=3, by=0.5)
s1
# [1] -2.0 -1.5 -1.0 -0.5 0.0 0.5 1.0 1.5 2.0
# [10] 2.5 3.0
```

Vector operations. You can also do calculations on vectors, for example using `x1` from above:

```
x1 * 2
# [1] 52 20 8 14 82 38
round(sqrt(x1*2.6), 2)
# [1] 8.22 5.10 3.22 4.27 10.32 7.03
```

Missing Values. Missing values are coded as `NA` in R. For example,

```
x2 <- c(3, -7, NA, 5, 1, 1)
x2
# [1] 3 -7 NA 5 1 1
x3 <- c("Rat", NA, "Mouse", "Hamster")
x3
# [1] "Rat" NA "Mouse" "Hamster"
```

Managing Objects. Use function `ls()` to list the objects in your workspace. The `rm()` function removes (deletes) them.

```
ls()
# [1] "altura" "anchura" "area" "Ciudad"
# [5] "edad" "Edad" "EdaD" "longitud"
# [9] "Nombre" "op" "r1" "s1"
# [13] "saludo" "x" "x1" "x2"
# [17] "x3" "xm"
rm(x1, x2, x3, r1, s1, AgE, age)
```

```
# Warning in rm(x1, x2, x3, r1, s1, AgE, age):
# object 'AgE' not found
```

```
# Warning in rm(x1, x2, x3, r1, s1, AgE, age):
# object 'age' not found
```

```
ls()
# [1] "altura" "anchura" "area" "Ciudad"
# [5] "edad" "Edad" "EdaD" "longitud"
# [9] "Nombre" "op" "saludo" "x"
# [13] "xm"
```

Exercise. Calculate the gross by adding the tax to net amount.

```
net <- c(108.99, 291.42, 16.28, 62.29, 31.77)
tax <- c(22.89, 17.49, 0.98, 13.08, 6.67)
```

7. R Functions and Packages

R Functions. We have already used some R functions (e.g. `c()`, `mean()`, `rep()`, `sqrt()`, `round()`). Most of the computations in R involves using functions. A function essentially has a name and a list of arguments separated by a comma. Let's have look at an example:

```
seq(from = 5, to = 8, by = 0.4)
# [1] 5.0 5.4 5.8 6.2 6.6 7.0 7.4 7.8
```

The function name is `seq` and it has three arguments `from`, `to` and `by`. The arguments `from` and `to` are the start and end values of a sequence that you want to create, and `by` is the increment of the sequence. The `seq()` functions has other arguments that you could use which are documented in the help page. For example, we could use the argument `length.out` (instead of `by`) to fix the length of the sequence as follows:

```
seq(from = 5, to = 8, length.out = 16)
# [1] 5.0 5.2 5.4 5.6 5.8 6.0 6.2 6.4 6.6 6.8 7.0
# [12] 7.2 7.4 7.6 7.8 8.0
```

Custom Functions. You can create your own functions (using the `function()` keyword) which is a very powerful way to extend R. Writing your own functions is outside the scope of this guide. As you get more and more familiar with R it is very likely that you will need to learn how to do so but for now you don't need to.

R Packages. You can already do many things with a standard R installation—but it can be extended using contributed packages. Packages are like apps for R. They can contain functions, data and documentation.

Extending Base R. Base R already comes with over two-thousand functions that have been proven to be versatile, reliable and stable. That is no small feat. When it is possible to solve a problem with fewer external dependencies, doing so follows time-honoured best practices. You want to think carefully before adding dependencies.

The tinyverse View. The philosophy of *less is more* is at the core of the tinyverse⁴. Fewer dependencies means a smaller footprint, faster installation, and most importantly fewer nodes in your dependency graph. Experience, as well as empirical and theoretical software engineering practice have demonstrated that failure increases with complexity.

⁴Tinyverse: <http://www.tinyverse.org/>

So choosing when to rely on additional packages has to balance the increased functionality a package brings with both its history of development, its development model, maintenance status, and history of both changes and fixes. This is a complex topic, and there are no easy answers. But by adding another package, we always open a door to interface changes we no longer control. The added functionality is clearly valuable at times, yet one has to remain aware of the costs that may accrue as a consequence. So this document takes the view that *fewer is better*, and will rely on only two additional packages: `data.table`⁵ for data wrangling as well as input/output, and `ggplot2`⁶ for visualization.

Installation. If needed, install these two packages via the following command which should pick the suitable version for your installation:

```
install.packages(c("data.table", "ggplot2"))
```

8. Chick Weight Data

R comes with many datasets installed⁷. We will use the `ChickWeight` dataset to learn about data manipulation. The help system gives a basic summary of the experiment from which the data was collect:

"The body weights of the chicks were measured at birth and every second day thereafter until day 20. They were also measured on day 21. There were four groups of chicks on different protein diets."

You can get more information, including references by typing:

```
help("ChickWeight")
```

The Data. There are 578 observations (rows) and 4 variables:

- Chick – unique ID for each chick.
- Diet – one of four protein diets.
- Time – number of days since birth.
- weight – body weight of chick in grams.

Note. `weight` has a lower case `w` (recall R is case sensitive).

Objective. Investigate the *effect of diet on the weight over time*.

9. Importing The Data

First we will import the data from a file called `ChickWeight.csv` using the `fread()` function from the `data.table` package which returns a `data.table` object (whereas the dataset built into R has a different format). The first thing to do, outside of R, is to open the file `ChickWeight.csv` to check what it contains and that it makes sense. Now we can import the data as follows:

```
suppressMessages(library(data.table)) # tinyverse
cw <- fread("ChickWeight.csv")
```

⁵data.table: <http://r-datatable.com>

⁶ggplot2: <https://ggplot2.tidyverse.org>

⁷Telee data() en la consola de R para ver una lista de todos los datasets.

Important Note. If all goes well then the data is now stored in an R object called `cw`. If you get the following error message then you need to change the working directory to where the data is stored.

Error: 'ChickWeight.csv' does not exist in current working directory ...

Change the working directory in RStudio. From the menu bar select "Session - Set Working Directory - Choose Directory..." then go to the directory where the data is stored. Alternatively, within in R, you could use the function `setwd()`⁸. You can also specify a full path, using `~` to denote your home directory.

10. Looking at the Dataset

To look at the data type just type the object (dataset) name:

```
cw
#      Chick Diet Time weight
#      1:      1   1    0     42
#      2:      1   1    2     51
#      3:      1   1    4     59
#      ---
# 576:     50    4   18    234
# 577:     50    4   20    264
# 578:     50    4   21    264
```

Several base R functions help us inspect the data: `str()` compactly displays the structure, `summary()` provides a summary, and `head()` and `tail()` display the beginning and end of the data set.

```
str(cw)
# Classes 'data.table' and 'data.frame':
#      578 obs. of 4 variables:
# $ Chick : int [1:578] 1 1 1 1 1 ...
# $ Diet  : int [1:578] 1 1 1 1 1 ...
# $ Time  : int [1:578] 0 2 4 6 8 ...
# $ weight: int [1:578] 42 51 59 64 76 ...
# - attr(*, ".internal.selfref")=<externalptr>
summary(cw)
#      Chick      Diet
#      Min.   : 1.0   Min.   :1.00
#      1st Qu.:13.0   1st Qu.:1.00
#      Median :26.0   Median :2.00
#      Mean   :25.8   Mean   :2.24
#      3rd Qu.:38.0   3rd Qu.:3.00
#      Max.   :50.0   Max.   :4.00
#      Time      weight
#      Min.   : 0.0   Min.   : 35
#      1st Qu.: 4.0   1st Qu.: 63
#      Median :10.0   Median :103
#      Mean   :10.7   Mean   :122
#      3rd Qu.:16.0   3rd Qu.:164
#      Max.   :21.0   Max.   :373
```

Interpretation. This shows that the dataset has 578 observations and 4 variables as we would expect, and as compared to the original data file `ChickWeight.csv`. So a good start. `str()` call notes the types of variables (all integer here) and the first few values. The RStudio 'Environment' pane provides a very similar view.

⁸Use `getwd()` para el directorio actual y `setwd("/to/data/path/data.csv")` para cambiarlo.


Exercise. It is important to look at the last observations of the dataset as it could reveal potential data issues. Use the `tail()` function to do this. Is it consistent with the original data file `ChickWeight.csv`?

11. Chick Weight: Data Visualisation

ggplot2 Package. To visualise the chick weight data, we will use the `ggplot2` package. Our interest is in seeing how the weight changes over time for the chicks by diet. For the moment don't worry too much about the details just try to build your own understanding and logic. To learn more try different things even if you get an error messages.

First plot. Let's plot the weight data (vertical axis) over time (horizontal axis).


```
# (Silently) load the plotting package
suppressMessages(library(ggplot2))
# An empty plot (the plot on the left)
ggplot(cw, aes(Time, weight))
# With data (the plot on the right)
ggplot(cw, aes(Time, weight)) + geom_point()
```



Exercise. Switch the variables Time and weight in code used for the plot on the right? What do you think of this new plot compared to the original?

Add colour for Diet. The graph above does not differentiate between the diets. Let's use a different colour for each diet.

```
# Adding colour for diet
ggplot(cw, aes(Time, weight, colour=factor(Diet))) +
  geom_point()
```



Interpretation. It is difficult to conclude anything from this graph as the points are printed on top of one another (with diet 1 underneath and diet 4 at the top).

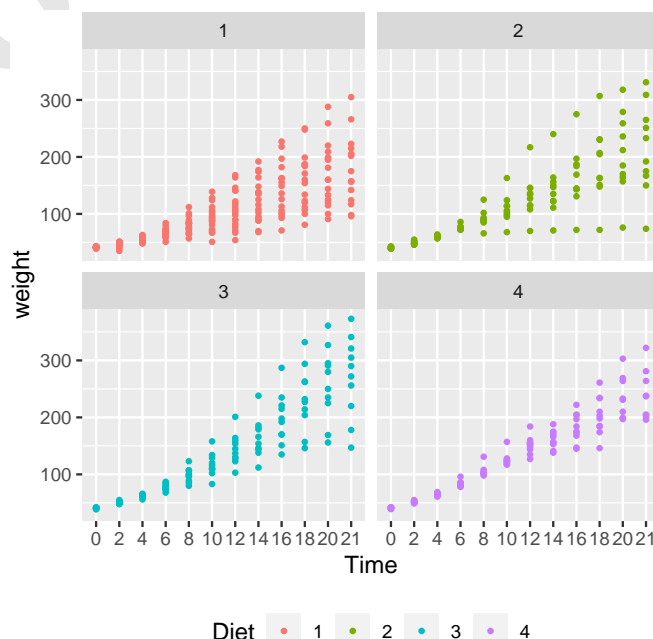
Factor Variables. Before we continue, we have to make an important change to the `cw` dataset by making Diet and Time factor variables. This means that R will treat them as categorical variables instead of continuous variables. It will simplify our coding.

```
cw[, Diet := factor(Diet)]
cw[, Time := factor(Time)]
str(cw) # notice the difference ?
# Classes 'data.table' and 'data.frame':
# 578 obs. of 4 variables:
# $ Chick : int [1:578] 1 1 1 1 1 ...
# $ Diet : Factor w/ 4 levels "1","2","3","4":
# 1 1 1 1 1 ...
# $ Time : Factor w/ 12 levels
# "0","2","4","6",...: 1 2 3 4 5 ...
# $ weight: int [1:578] 42 51 59 64 76 ...
# - attr(*, ".internal.selfref")=<externalptr>
```

Notice that the `:=` operator altered the variable “in-place”, and no explicit assignment was made. This is a key feature of `data.table` which operated “by reference”: changes are made in reference to one instance of the `cw` variable, rather than by creating updated copies. We will revisit this `:=` assignment below.

facet_wrap() function. To plot each diet separately in a grid using `facet_wrap()`:

```
# Adding jitter to the points
ggplot(cw, aes(Time, weight, colour=Diet)) +
  geom_point() +
  facet_wrap(~ Diet) +
  theme(legend.position = "bottom")
```



Exercise. To overcome the issue of overlapping points we can *jitter* the points using `geom_jitter()`. Replace the `geom_point()` above with `geom_jitter()`. What do you observe?


Interpretation. Diet 4 has the least variability but we can't really say anything about the mean effect of each diet although diet 3 seems to have the highest.

Exercise. For the `legend.position` try using “top”, “left” and “none”. Do we really need a legend for this plot?

Mean line plot. Next we will plot the mean changes over time for each diet using the `stat_summary()` function:

```
ggplot(cw, aes(Time, weight,
               group=Diet, colour=Diet)) +
  stat_summary(fun.y="mean", geom="line")
```

```
# Warning: `fun.y` is deprecated. Use `fun` instead.
```



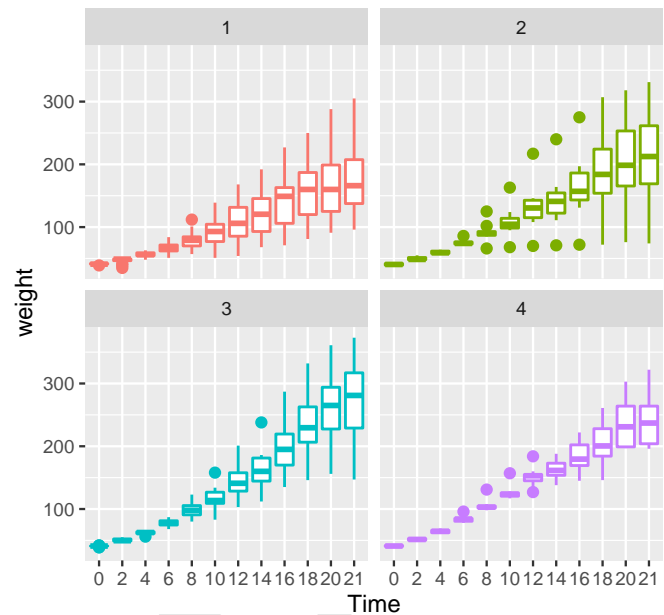
Interpretation. We can see that diet 3 has the highest mean weight gain by the end of the experiment but we don't have any information about the variation (uncertainty) in the data.

Exercise. What happens when you add `geom_point()` to the plot above? Don't forget the `+`. Does it make a difference if you put it before or after the `stat_summary(...)` line? Hint: Look very carefully at how the graph is plotted.

Box-whisker plot. To see variation between the different diets we use `geom_boxplot` to plot a box-whisker plot. A note of caution is that the number of chicks per diet is relatively low to produce this plot.

```
ggplot(cw, aes(Time, weight, colour=Diet)) +
  facet_wrap(~ Diet) +
  geom_boxplot() +
  theme(legend.position = "none") +
  ggtitle("Chick Weight over Time by Diet")
```

Chick Weight over Time by Diet



Interpretation. Diet 3 seems to have the highest “average” weight gain but it has more variation than diet 4 which is consistent with our findings so far.

Exercise. Add the following information to the above plot:


- x-axis label (use `xlab()`): “Time (days)”
- y-axis label (use `ylab()`): “Weight (grams)”

Final Plot. Let's finish with a plot that you might include in a publication.

```
ggplot(cw, aes(Time, weight, group=Diet,
               colour=Diet)) +
  facet_wrap(~ Diet) +
  geom_jitter() +
  stat_summary(fun.y="mean", geom="line",
               colour="black") +
  theme(legend.position = "none") +
  ggtitle("Chick Weight over Time by Diet") +
  xlab("Time (days)") +
  ylab("Weight (grams)")
```

```
# Warning: `fun.y` is deprecated. Use `fun` instead.
```

Chick Weight over Time by Diet



12. data.table Data Wrangling Basics

In this section we will learn how to wrangle (manipulate) datasets using the `data.table` package. Conceptually, `data.table` operations can be viewed as `dt[i, j, by]` with some intentional similarity to SQL. Here `i` can select (or subset) rows, `j` is used to select, summarise or mutate columns, and `by` is the grouping operator. Numerous examples follow.

j to select (or transform) columns. Adds a new variable (column) or modifies an existing one. We already used this above to create factor variables.

```
cw[, weightKg := weight/1000] # add a column
cw
#      Chick Diet Time weight weightKg
# 1:      1    1    0     42     0.042
# 2:      1    1    2     51     0.051
# 3:      1    1    4     59     0.059
# ---
# 576:    50    4   18    234     0.234
# 577:    50    4   20    264     0.264
# 578:    50    4   21    264     0.264
cw[, Diet := paste0("Diet_", Diet)] # mod col.
cw
#      Chick  Diet Time weight weightKg
# 1:      1 Diet_1    0     42     0.042
# 2:      1 Diet_1    2     51     0.051
# 3:      1 Diet_1    4     59     0.059
# ---
# 576:    50 Diet_4   18    234     0.234
# 577:    50 Diet_4   20    264     0.264
# 578:    50 Diet_4   21    264     0.264
```

j to select (or transform) columns. Keeps, drops or reorders variables.

```
# Keep variables Time, Diet and weightKg
cw[, .(Chick, Time, Diet, weightKg)]
#      Chick Time  Diet weightKg
# 1:      1    0 Diet_1    0.042
# 2:      1    2 Diet_1    0.051
# 3:      1    4 Diet_1    0.059
# ---
# 576:    50   18 Diet_4    0.234
# 577:    50   20 Diet_4    0.264
# 578:    50   21 Diet_4    0.264
```

j to summarise. It can be used to create aggregations, which is particularly handy with the grouping operator. The following example computes means and standard deviations of the 'weight' variable grouped by 'Diet'. Note that the output has been truncated.

```
cw[, .(Mean=mean(weight), SDev=sd(weight)),
  by=(Diet, Time)]
#      Diet Time      Mean      SDev
# 1: Diet_1    0  41.4000  0.994723
# 2: Diet_1    2  47.2500  4.278157
# 3: Diet_1    4  56.4737  4.128067
# ---
# 46: Diet_4   18 202.9000 33.557413
# 47: Diet_4   20 233.8889 37.568086
# 48: Diet_4   21 238.5556 43.347754
```

setnames() to name or rename. Renames variables whilst keeping all variables.

```
setnames(cw, c("Diet", "weight"),
  c("Group", "Weight"))
cw
#      Chick Group Time Weight weightKg
# 1:      1 Diet_1    0     42     0.042
# 2:      1 Diet_1    2     51     0.051
# 3:      1 Diet_1    4     59     0.059
# ---
# 576:    50 Diet_4   18    234     0.234
# 577:    50 Diet_4   20    264     0.264
# 578:    50 Diet_4   21    264     0.264
```

i operator. Keeps or drops observations (rows).

```
cw[Time == 21 & Weight > 300]
#      Chick Group Time Weight weightKg
# 1:      7 Diet_1   21    305     0.305
# 2:     21 Diet_2   21    331     0.331
# 3:     29 Diet_2   21    309     0.309
# 4:     32 Diet_3   21    305     0.305
# 5:     34 Diet_3   21    341     0.341
# 6:     35 Diet_3   21    373     0.373
# 7:     40 Diet_3   21    321     0.321
# 8:     48 Diet_4   21    322     0.322
```

For comparing values in vectors use: `<` (less than), `>` (greater than), `<=` (less than and equal to), `>=` (greater than and equal to), `==` (equal to) and `!=` (not equal to). These can be combined logically using `&` (and) and `|` (or).

Keying observations. Setting a key changes the order of the observations (rows), and also makes indexing faster.

```

cw[order(Weight)]      # on the fly
#      Chick  Group Time Weight weightKg
#    1:    18 Diet_1  2     35    0.035
#    2:     3 Diet_1  2     39    0.039
#    3:    18 Diet_1  0     39    0.039
#    ---
# 576:    34 Diet_3  21    341    0.341
# 577:    35 Diet_3  20    361    0.361
# 578:    35 Diet_3  21    373    0.373
setkey(cw, Chick, Time) # setting a key
cw
#      Chick  Group Time Weight weightKg
#    1:     1 Diet_1  0     42    0.042
#    2:     1 Diet_1  2     51    0.051
#    3:     1 Diet_1  4     59    0.059
#    ---
# 576:    50 Diet_4  18    234    0.234
# 577:    50 Diet_4  20    264    0.264
# 578:    50 Diet_4  21    264    0.264

```

Exercise. What does the `order()` do? Try using `order(Time)` and `order(-Time)` in the `i` column.

13. Chaining

You may want to do multiple data wrangling steps at once. This is where the ‘chaining’ of `data.table` operations (i.e., several sets of commands with square brackets) comes to the rescue:

```

cw21 <- cw[Time %in% c(0,21)][      # i: select rows
, weight := Weight][              # j: mutate
, Group := factor(Group)][
, .(Chick, Group, Time, weight)][   # j: arrange
order(Chick, Time)][              # i: order
1:5]                               # i: subset

```

14. Chick Weight: Summary Statistics

From the data visualisations above we concluded that the diet 3 has the highest mean and diet 4 the least variation. In this section, we will quantify the effects of the diets using summary statistics. We start by looking at the number of observations and the mean of weight grouped by **diet** and **time**.

```

cw[, .(N      = .N,                # .N is nb per group
      Mean = mean(Weight)),        # compute mean
by=.(Group, Time)][              # group by Diet + Time
1:5]                             # display rows 1 to 5
#      Group Time  N      Mean
# 1: Diet_1    0  20  41.4000
# 2: Diet_1    2  20  47.2500
# 3: Diet_1    4  19  56.4737
# 4: Diet_1    6  19  66.7895
# 5: Diet_1    8  19  79.6842

```

by= argument. For each distinct combination of Diet and Time, the chick weight data is summarised into the number of observations (N, using the internal variable `.N` denoting current group size) and the mean (Mean) of weight.

Other summaries. We can calculate the standard deviation, median, minimum and maximum values—only at days 0 and 21.

```

cws <- cw[Time %in% c(0,21),
.(N      = .N,
  Mean   = mean(Weight),
  SDev   = sd(Weight),
  Median = median(Weight),
  Min    = min(Weight),
  Max    = max(Weight) ),
by=.(Group, Time)]

cws
#      Group Time  N      Mean   SDev Median Min Max
# 1: Diet_1    0  20  41.4   0.995  41.0  39  43
# 2: Diet_1    2  16 177.8  58.702 166.0  96 305
# 3: Diet_2    0  10  40.7   1.494  40.5  39  43
# 4: Diet_2    2  10 214.7  78.138 212.5  74 331
# 5: Diet_3    0  10  40.8   1.033  41.0  39  42
# 6: Diet_3    2  10 270.3  71.623 281.0 147 373
# 7: Diet_4    0  10  41.0   1.054  41.0  39  42
# 8: Diet_4    2   9 238.6  43.348 237.0 196 322

```

Finally, we can make the summaries “prettier” for a possible report or publication where we format the numeric values as text.

```

cws[, Mean_SD := paste0(format(Mean,digits=1),
" (",
format(SDev,digits=2),
")")]
cws[, Range := paste(Min, "-", Max)]
prettySum <- cws[, .(Group, Time, N, Mean_SD,
Median, Range)][
order(Group, Time)]

prettySum
#      Group Time  N      Mean_SD Median      Range
# 1: Diet_1    0  20  41 ( 0.99)  41.0  39 - 43
# 2: Diet_1    2  16 178 (58.70) 166.0  96 - 305
# 3: Diet_2    0  10  41 ( 1.49)  40.5  39 - 43
# 4: Diet_2    2  10 215 (78.14) 212.5  74 - 331
# 5: Diet_3    0  10  41 ( 1.03)  41.0  39 - 42
# 6: Diet_3    2  10 270 (71.62) 281.0 147 - 373
# 7: Diet_4    0  10  41 ( 1.05)  41.0  39 - 42
# 8: Diet_4    2   9 239 (43.35) 237.0 196 - 322

```

Final Table. Eventually you should be able to produce a publication-ready version such as the following table. Its code uses the `kable` and `kableExtra` packages. While the code is not displayed here for compactness, full details are of course in the sources.

| Group | Time | N | Mean_SD | Median | Range |
|--------|------|----|-------------|--------|-----------|
| Diet_1 | 0 | 20 | 41 (0.99) | 41.0 | 39 - 43 |
| Diet_1 | 21 | 16 | 178 (58.70) | 166.0 | 96 - 305 |
| Diet_2 | 0 | 10 | 41 (1.49) | 40.5 | 39 - 43 |
| Diet_2 | 21 | 10 | 215 (78.14) | 212.5 | 74 - 331 |
| Diet_3 | 0 | 10 | 41 (1.03) | 41.0 | 39 - 42 |
| Diet_3 | 21 | 10 | 270 (71.62) | 281.0 | 147 - 373 |
| Diet_4 | 0 | 10 | 41 (1.05) | 41.0 | 39 - 42 |
| Diet_4 | 21 | 9 | 239 (43.35) | 237.0 | 196 - 322 |

Interpretation. This summary table offers the same interpretation as before, namely that diet 3 has the highest mean and median

weights at day 21 but a higher variation than group 4. However it should be noted that at day 21, diet 1 lost 4 chicks from 20 that started and diet 4 lost 1 from 10. This could be a sign of some issues (e.g. safety).

Limitations of data. Information on bias reduction measures is not given and is not available either⁹. We don't know if the chicks were fairly and appropriately randomised to the diets and whether the groups are comparable (e.g., same breed of chicks, sex (gender) balance). Hence we should be very cautious with drawing conclusion and taking actions with this data.

15. Conclusion

This “Getting Started in R” guide introduced you to some of the basic concepts underlying R and used a real life dataset to produce some graphs and summary statistics. It is only a flavour of what R can do but hopefully you have seen some of power of R and its potential.

What next. There are plenty of R courses, books and on-line resources that you can learn from. It is hard to recommend any in particular as it depends on how you learn best. Find things that work for you (paying attention to the quality) and don't be afraid to make mistakes or ask questions. Most importantly have fun.

16. Acknowledgements

Special thanks to [Saghir Bashir](#) for publishing the initial version of *Getting Started with R*, [Brodie Gaslam](#) and [Matt Dowle](#) for feedback on this version, and to [Joshua Ulrich](#) for many discussions about *The Tinyverse*.

⁹I (ie Saghir) contacted the source authors and kindly received the following reply “*They were mainly undergraduate projects, final-year, rather than theses, so, unfortunately, it’s unlikely that any record remains, particularly after so many years.*”